

Machine Learning for Predicting Numeric Values

Nysret Musliu

Database and Artificial Intelligence Group (DBAI)

Example 1

- Predicting CPU performance

	Cycle time (ns)	Main memory (Kb)		Cache (Kb)	Channels		Performance
		MYCT	MMIN	MMAX	CACH	CHMIN	CHMAX
1	125	256	6000	256	16	128	198
2	29	8000	32000	32	8	32	269
...							
208	480	512	8000	32	0	0	67
209	480	1000	4000	0	0	0	45

Data set taken from Data Mining Practical Machine Learning Tools and Techniques by I. H. Witten, E. Frank, M. A. Hall and C. J. Pal

Example 2

.....

- Predicting house prices in Vienna

Nr.Rooms	Size (m ²)	Age of house	...	Price (K)
1	35	1		110
4	120	20		250
3	85	10		235
3	78	2		210
2	54	5		170
3	77	4		?
...				

Example 3

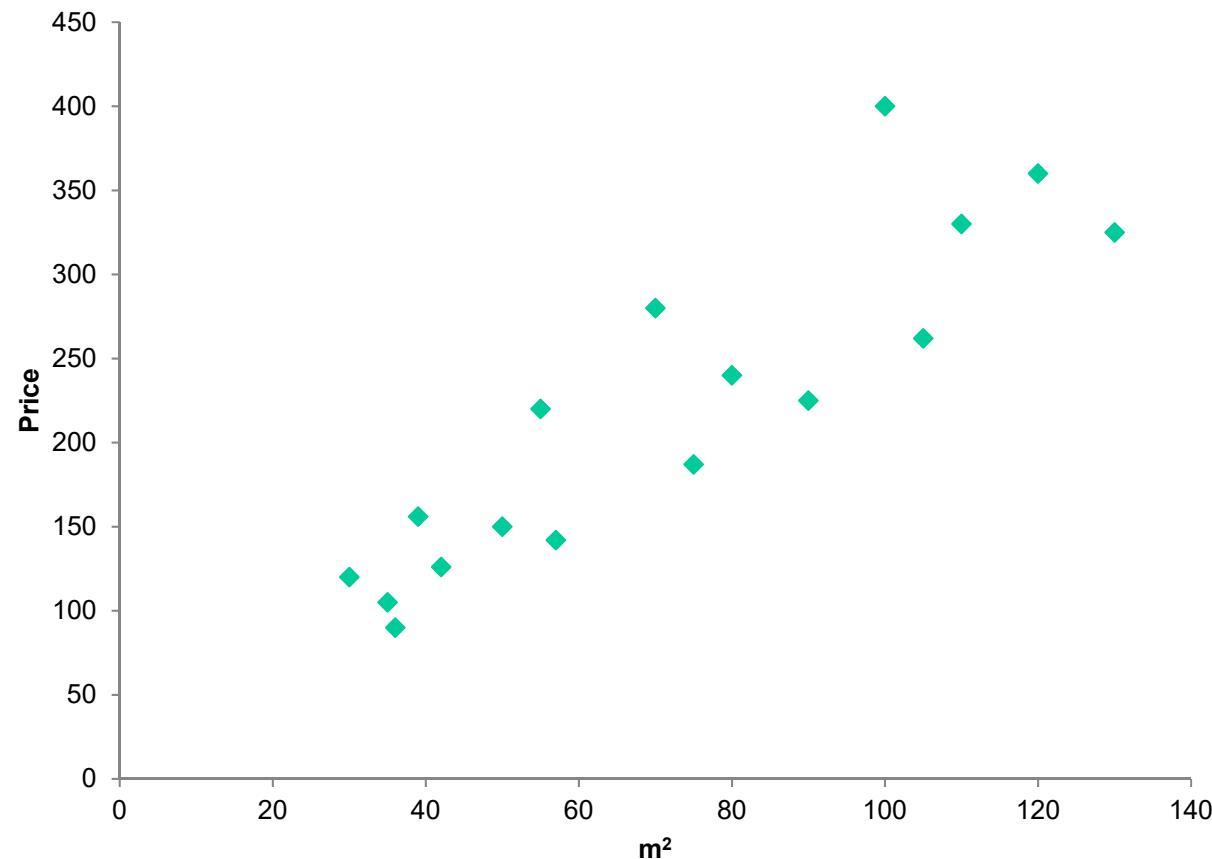
.....

- Predicting the running time of an algorithm (e.g. for solving of a SAT problem ... See SATzilla)

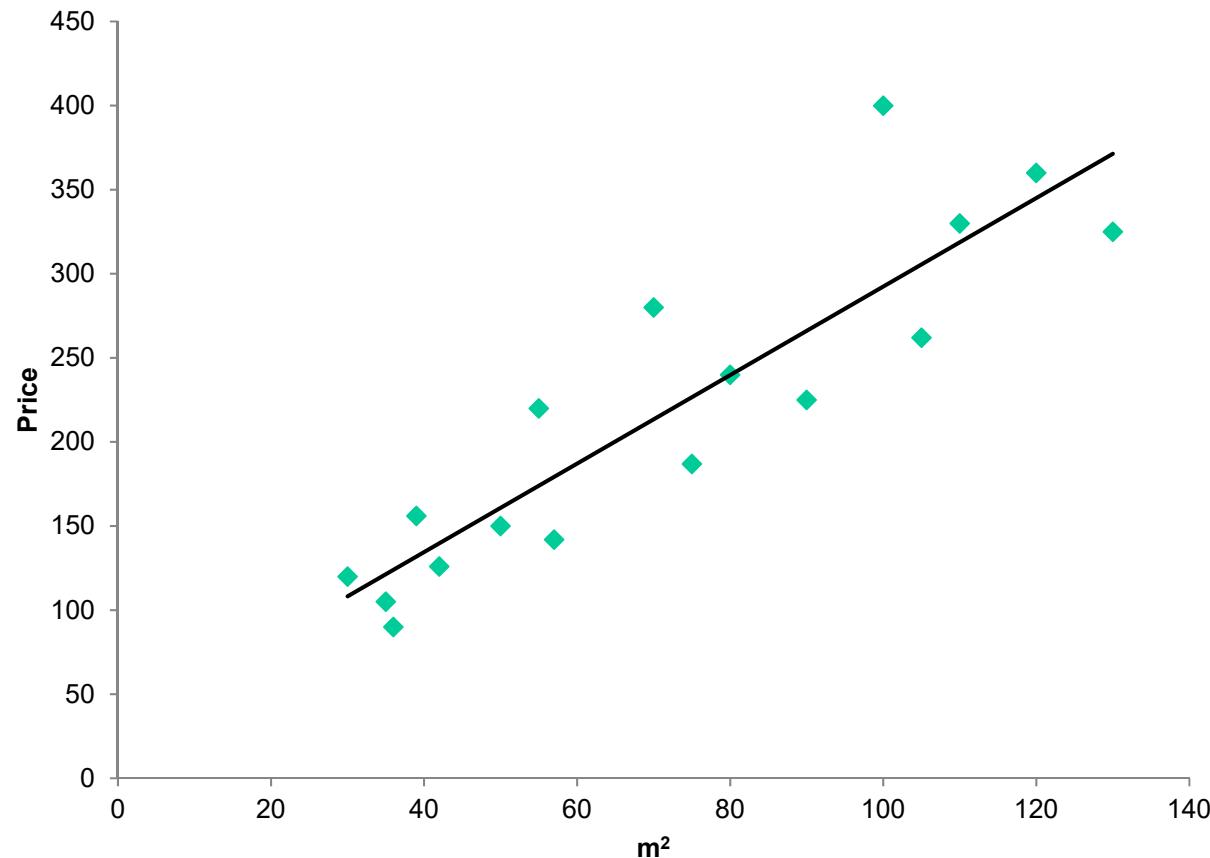
NrClauses	NrVariables	...	Time (sec)
100	80		10
4000	400		450
30000	8500		2350
300	78		25
2000	540		170
10000	450		?
...			

- Many other examples (see UCI repository, Kaggle...)

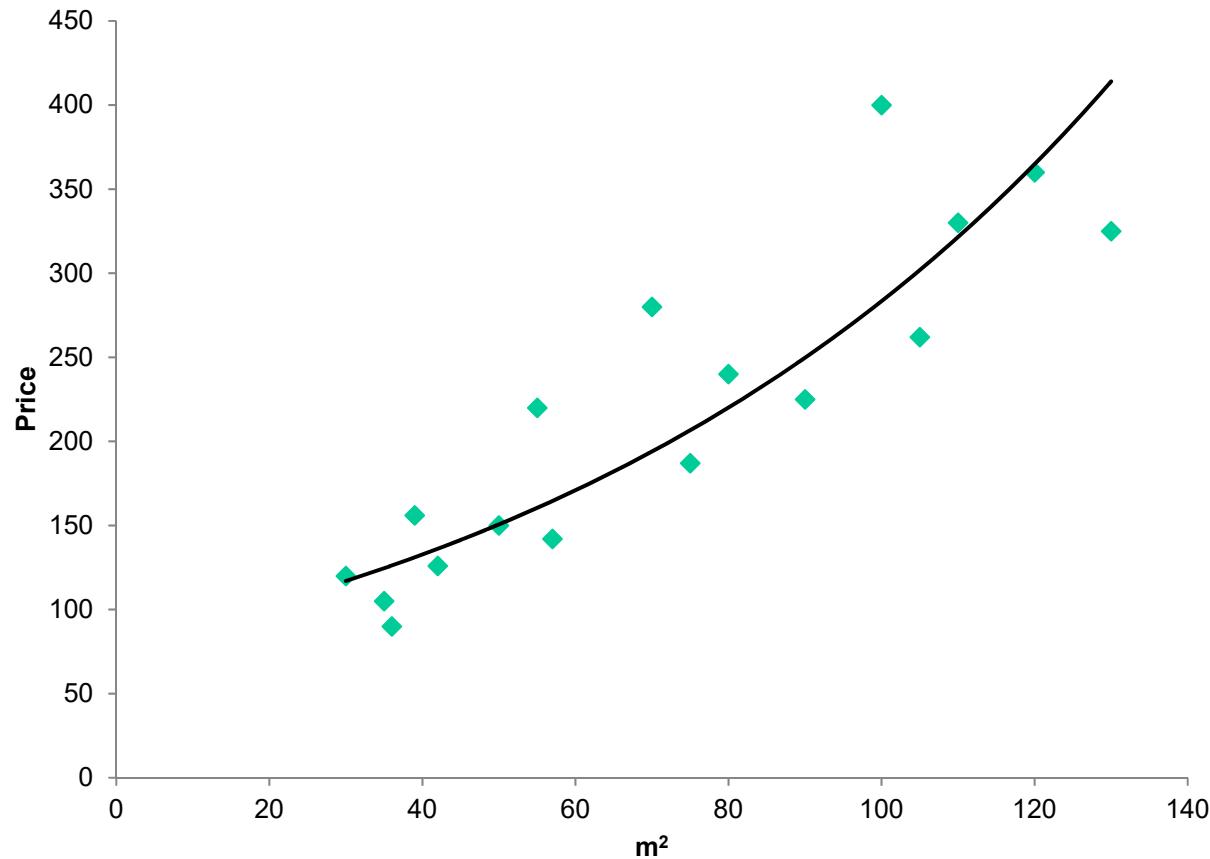
Prices vs. Square meter



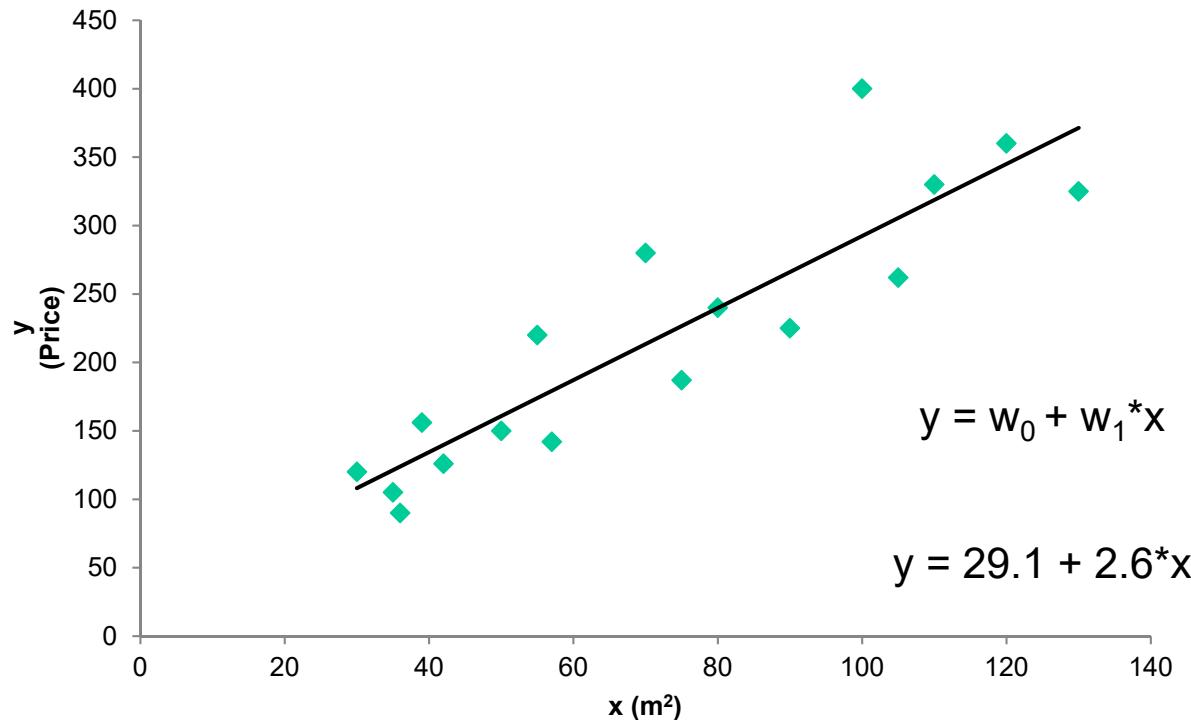
Model: Linear regression



Model: Other functions



Model: Linear regression



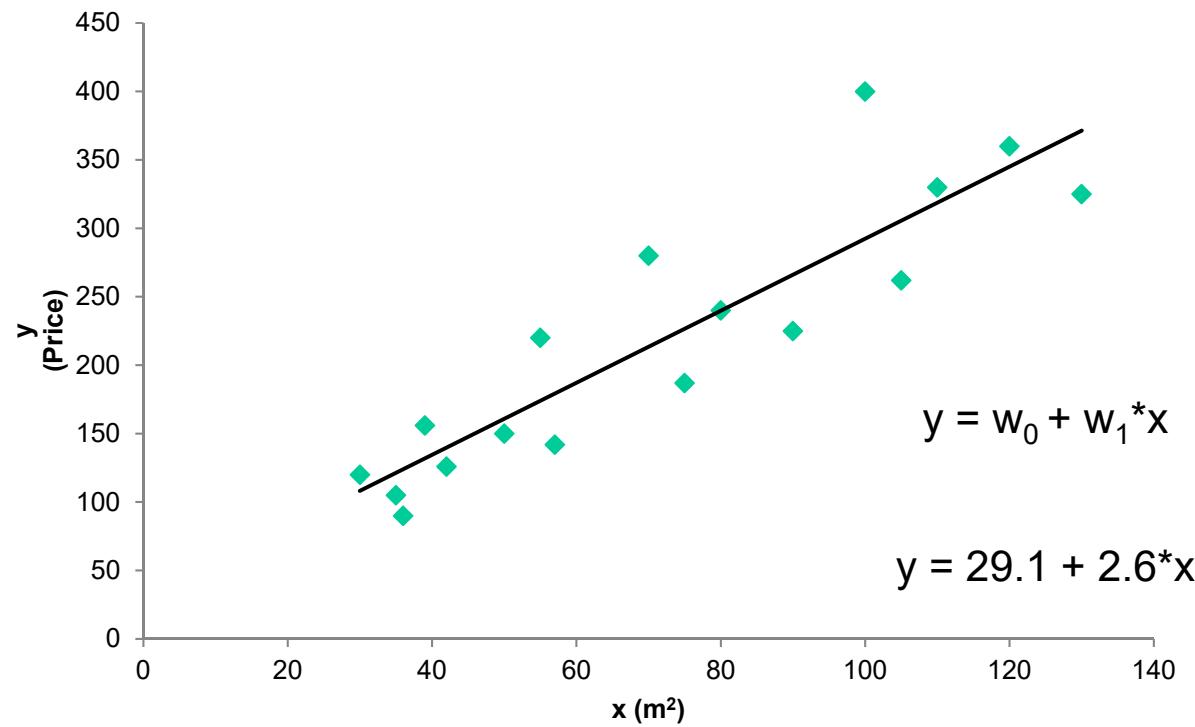
$w_0, w_1 \rightarrow$ parameters

Problem: Find best values for w_0, w_1

Minimize some cost metric (sum of squared errors of prediction, ...)

Residual sum of squares (RSS)

$$\text{RSS}(w_0, w_1) = \sum_{i=1}^N (y_i - (w_0 + w_1 * x_i))^2$$



Finding best parameters

Minimize cost over all possible w_0, w_1

$$\min \text{RSS}(w_0, w_1) = \sum_{i=1}^N (y_i - (w_0 + w_1 x_i))^2$$

Gradient descent algorithm

- Initial values for w_0, w_1
- Iterative change of these values
- Until convergence

Gradient descent

.....

Repeat until convergence

$$w_0 = w_0 - \alpha \frac{\partial}{\partial w_0} RSS(w_0, w_1)$$

$$w_1 = w_1 - \alpha \frac{\partial}{\partial w_1} RSS(w_0, w_1)$$

Update w_1 and w_0 simultaneously

α -> learning rate

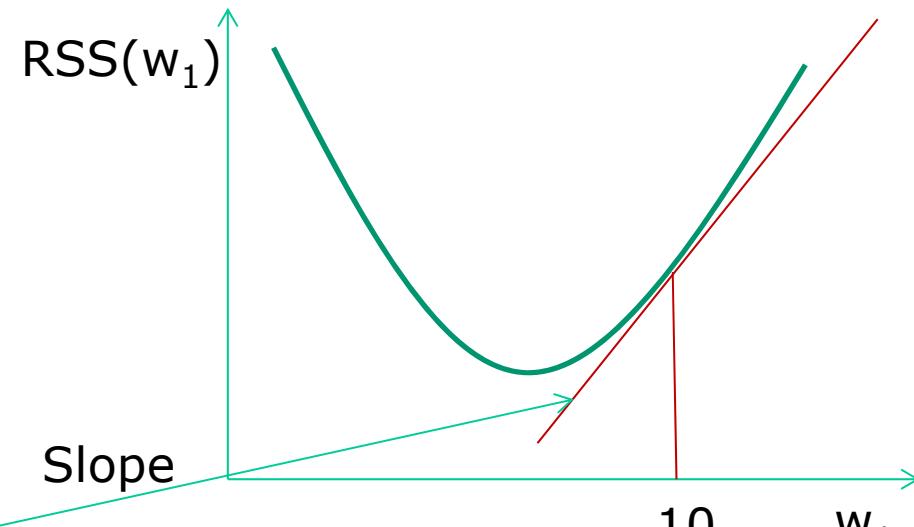
Gradient descent

Simple case: $y = w_1^*x$

Initial value $w_1 = 10$, $\alpha = 4$

Iteration 1:

$$w_1 = 10 - 4 * \frac{\partial}{\partial w_1} RSS(w_1)$$



Gradient descent

Simple case: $y = w_1^*x$

Initial value $w_1 = 10$, $\alpha = 4$

Iteration 1:

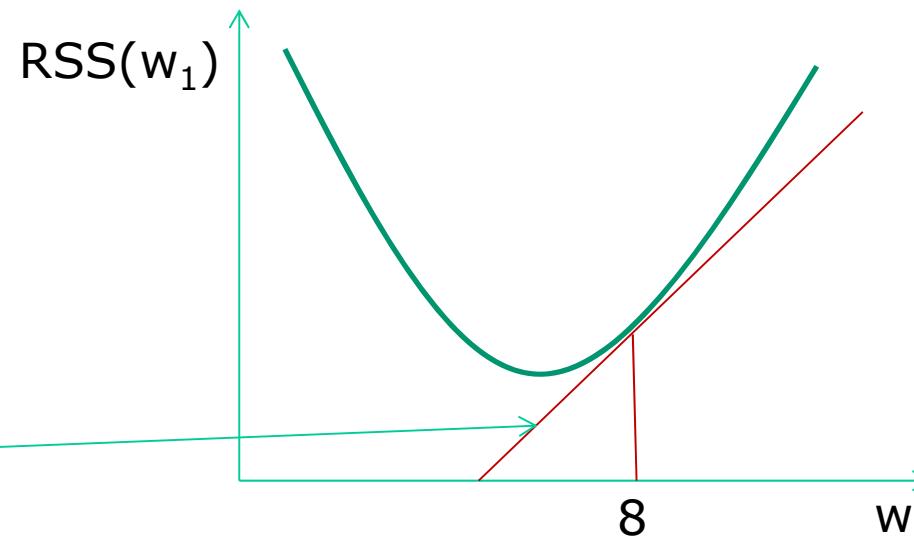
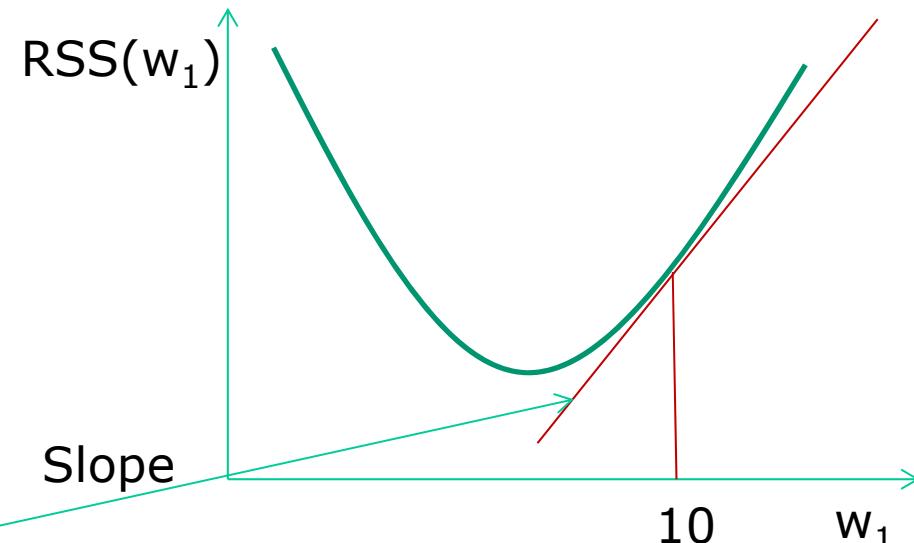
$$w_1 = 10 - 4 * \frac{\partial}{\partial w_1} RSS(w_1)$$

Suppose that this value is 0.5

Iteration 2

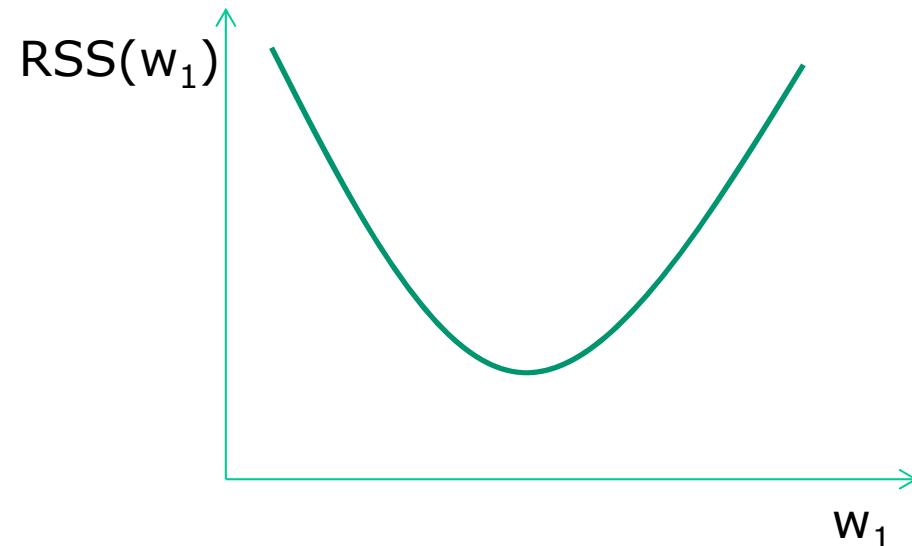
$$w_1 = 8 - 4 * \frac{\partial}{\partial w_1} RSS(w_1)$$

...

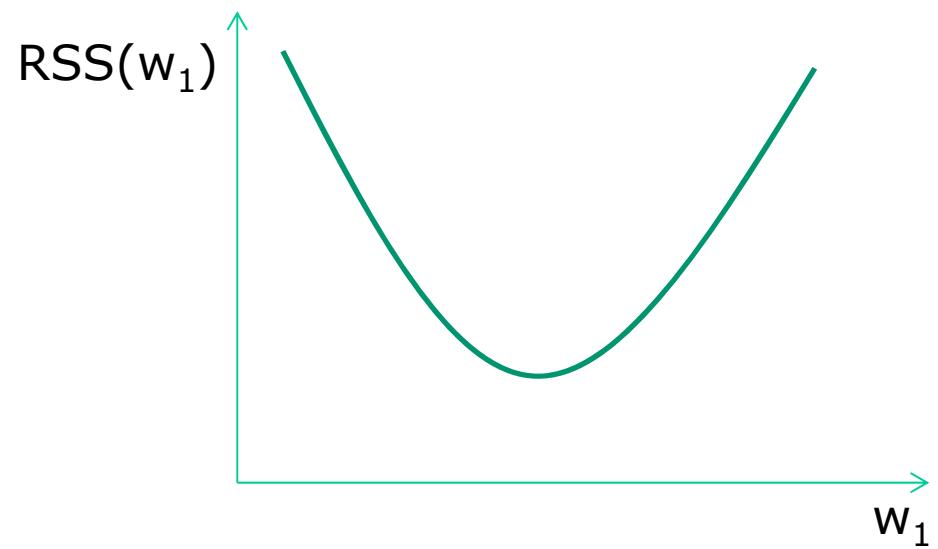


The impact of learning rate

Large α



Small α



Linear regression with multiple variables

.....

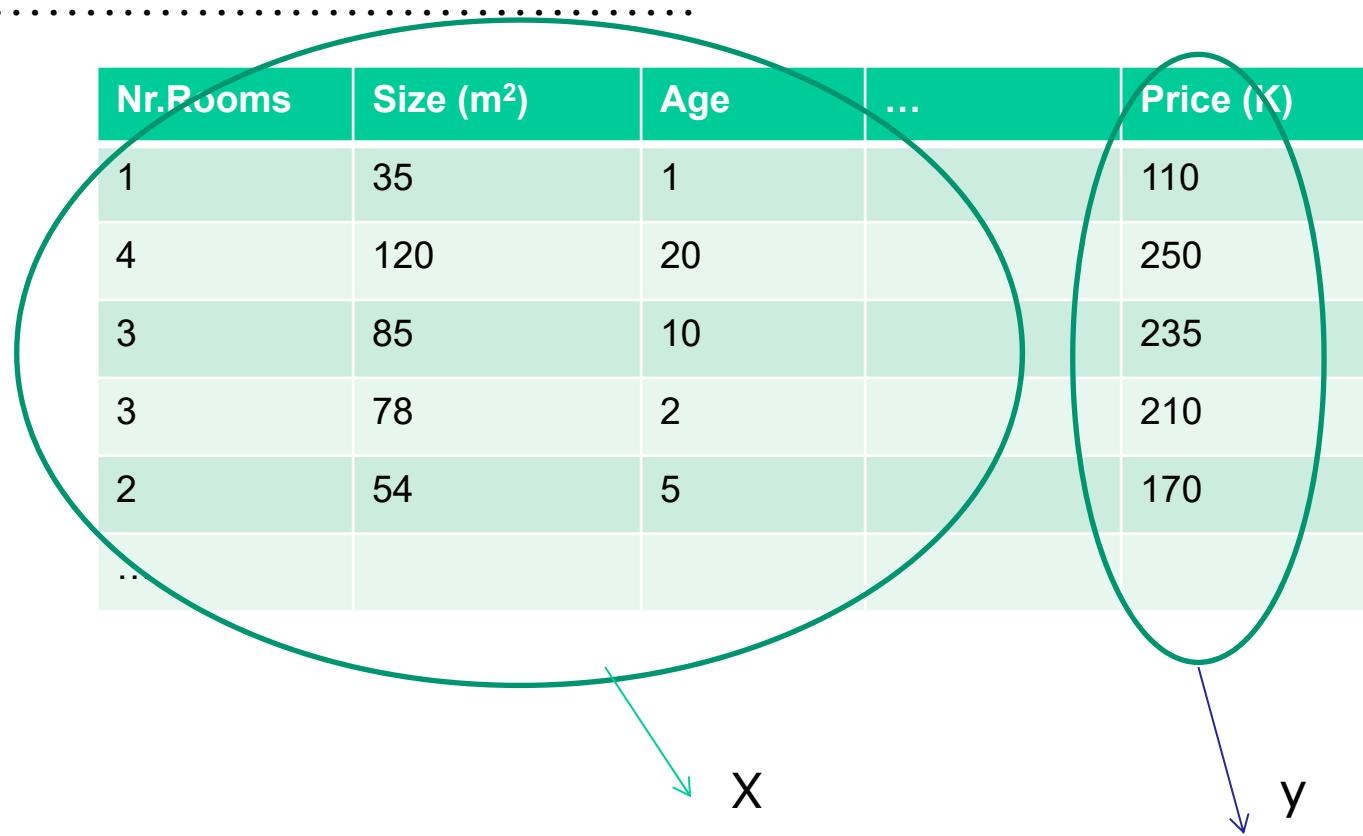
Gradient descent

Repeat until convergence

$$w_i = w_i - \alpha \frac{\partial}{\partial w_i} RSS(w_0, w_1, \dots, w_n)$$

Update w_0, w_1, \dots, w_n simultaneously

Analytical solution for linear regression

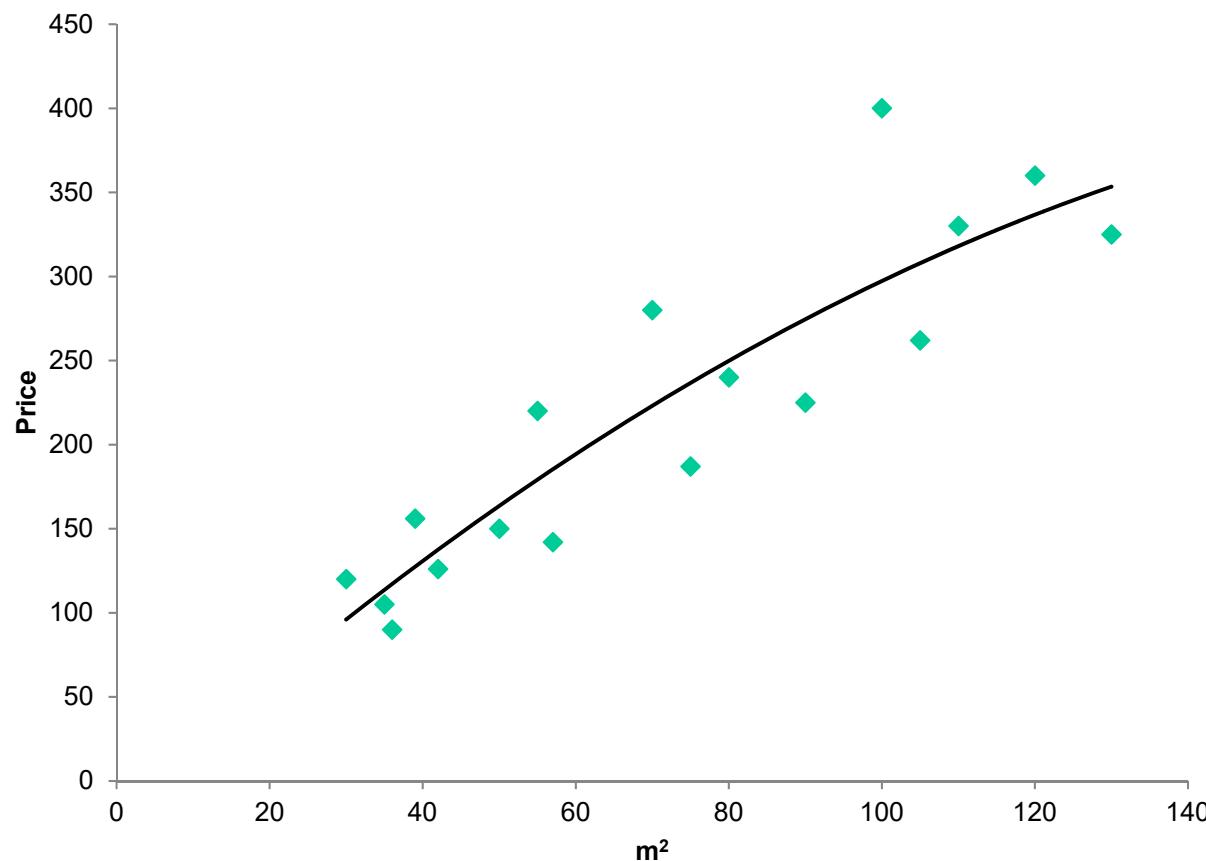


$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

Gradient descent versus analytical solution

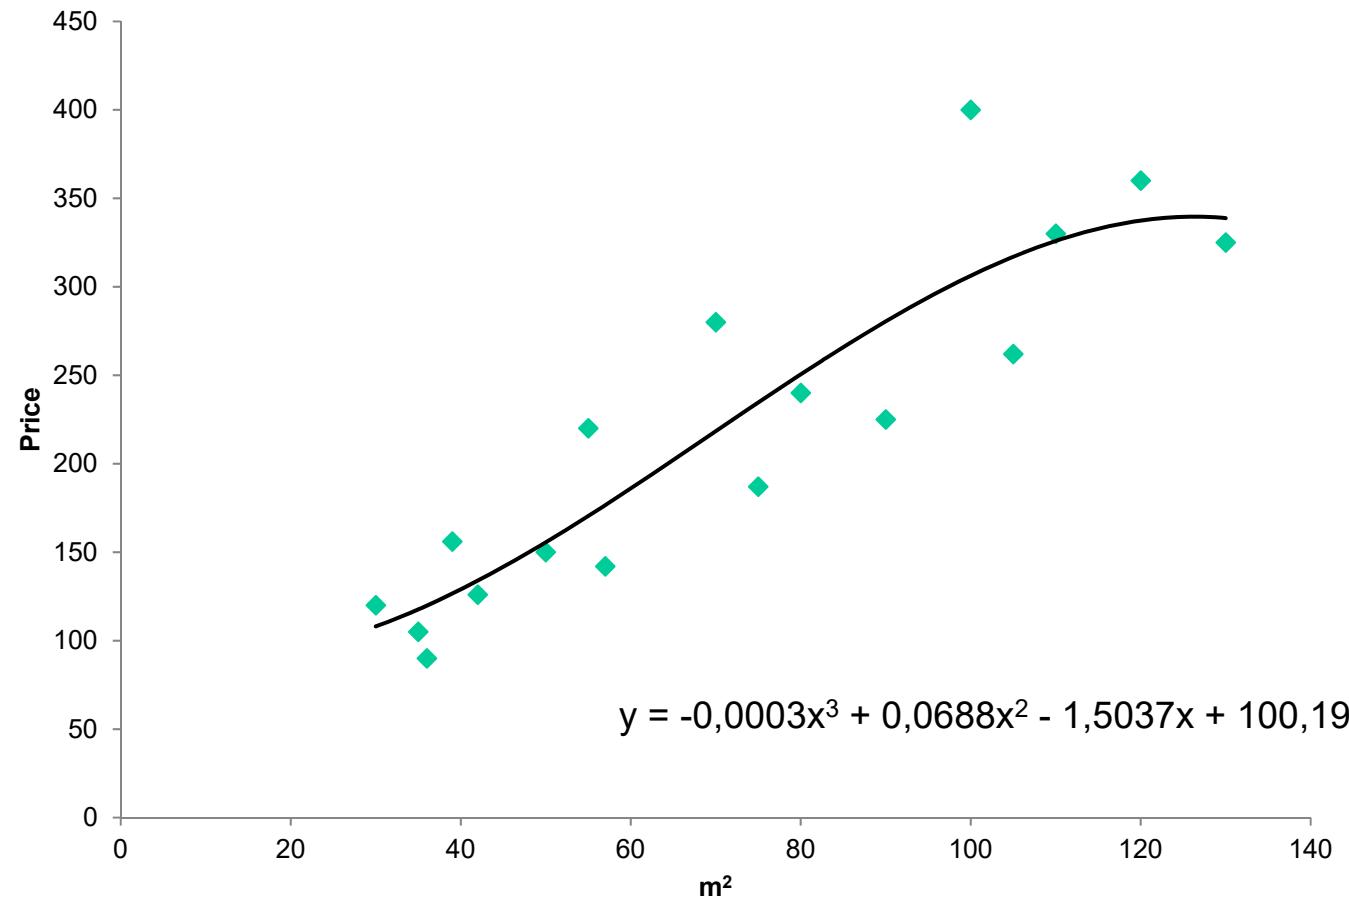
- Gradient Descent
 - Needs many iteration
 - Learning rate should be determined
 - Useful when the number of features is large
- Normal equation
 - No iterations
 - No need to choose learning rate
 - Works slow if the number of features is too large

Polynomial regression

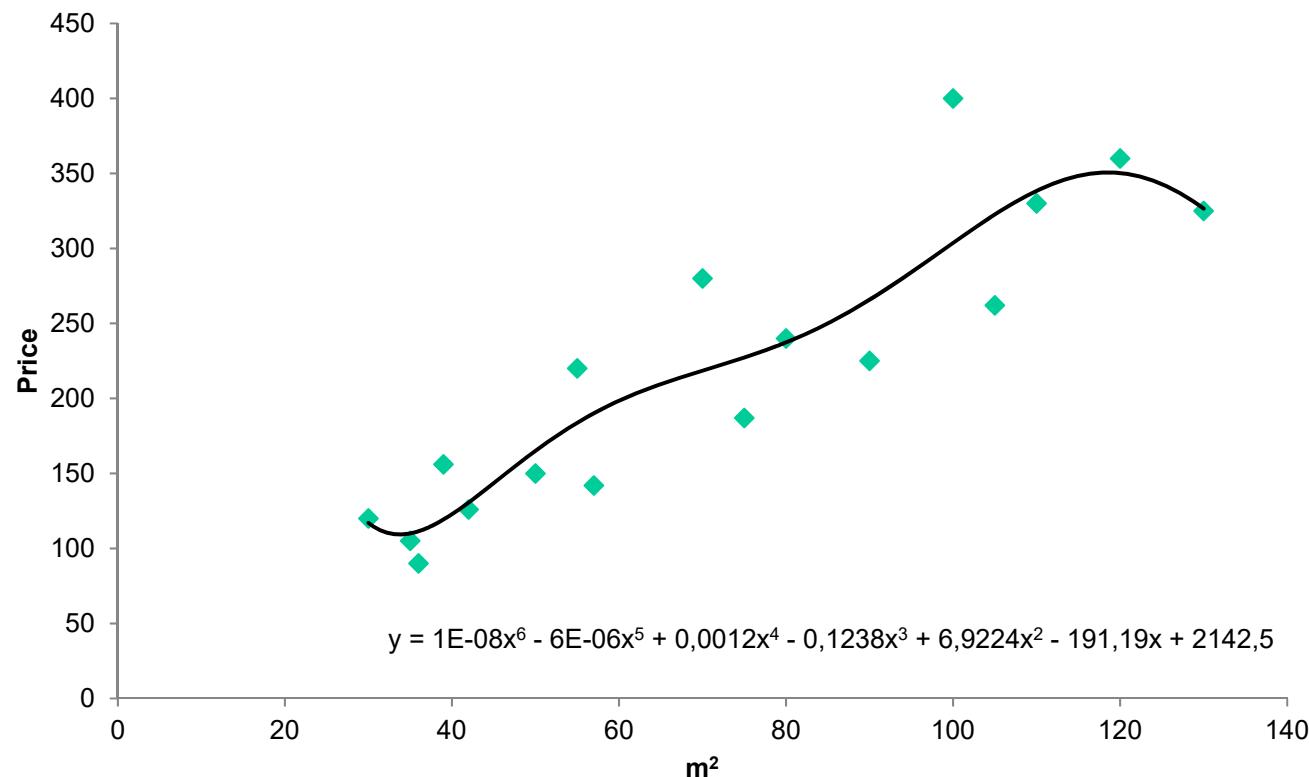


Polynomial regression

.....



Polynomial regression



- High-order polynomials
 - Flexible
 - Problem of overfitting
 - Usually coefficients get large

- Try to avoid overfitting caused by high-order polynomials
- Objective:
 - Balance between
 - How well function fits the data
 - Magnitude of coefficients

Total cost=measure of fit + $\lambda * \text{measure of magnitude of coefficients}$

Selection of λ

.....

- Large λ
 - High bias, low variance
- Small λ
 - Low bias, high variance
- λ controls the model complexity

- Use regularization to eliminate some features
- Some coefficients should get *exactly 0*
- Some features are not used
- Features are selected if the coefficients are non-zero

Ridge Regression (L_2 regularized regression).....

Total cost=measure of fit + $\lambda * \text{measure of magnitude of coefficients}$

$$\text{Total cost} = \text{RSS}(w_0, \dots, w_n) + \lambda * \|w\|_2^2$$

$$\|w\|_2^2 = w_0^2 + \dots + w_n^2$$

Encourages small weights but not exactly 0

Set small ridge coefficients to 0

Lasso Regression (L_1 regularized regression).....

Total cost=measure of fit + $\lambda * \text{measure of magnitude of coefficients}$

$$\text{Total cost} = \text{RSS}(w_0, \dots, w_n) + \lambda * \|w\|_1$$

$$\|w\|_1 = |w_0| + \dots + |w_n|$$

Lasso regression can be used for feature selection

- Feature selection task
 - Feature selection: select features that minimize redundancy and maximize the relevance to the target
- Advantages of feature selection
 - Improved learning performance
 - Lower computational complexity
 - Decreased storage

Feature Selection: all subsets

- Enumerate all possible subsets
 - Subsets for 8 features

1	0	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---

1	1	0	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---

1	0	1	0	0	0	0	0	0
---	---	---	---	---	---	---	---	---

1	0	0	1	0	0	0	0	0
---	---	---	---	---	---	---	---	---

1	0	0	0	1	0	0	0	0
---	---	---	---	---	---	---	---	---

...

1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---

Too many possibilities:

$2^{\text{Number of features}}$

- The size of search space for m features is $O(2^m)$
- NP-hard problem
- Wrapper Model: Search strategies
 - Best-first search
 - Hill Climbing
 - Genetic Algorithms
 - Branch and Bound
 - Tabu Search
 - ...

Greedy strategies

- Efficient and robust against overfitting
- Forward selection
 - Empty set of features
 - Features are added in the subset iteratively
- Backward elimination
 - Full set of features
 - Iterative elimination of less promising features

- Forward Search
 - $F=\{\}$
 - While the desired number of features arrived
 - For each feature f not in F
 - Estimate by cross-validation the error of model on $F \cup f$
 - Add f with lowest error to F
- Backward elimination
 - $F=\{\text{All features}\}$
 - While not reduced to desired number of features
 - For each feature f in F :
 - Estimate by cross-validation the error of model on F/f
 - Remove from F the feature f with the highest error

Other techniques

- Metahuristic based techniques
- Wrapper models give better results than filter models
- Computationally expensive
- Better performance for the predefined classifier

- More information will be given in the next lectures

- Data set of pairs $(x_1, y_1), (x_2, y_2), \dots, (x_n, y_n)$
- $y_1, y_2, \dots, y_n \rightarrow$ numeric values
- Prediction for the new instance x_m with k-NN:
 - Similar to k-NN for classification
 - Find k closest x_i in the data set: $x_{n1}, x_{n2}, \dots, x_{ni}$
 - Prediction: A simple implementation would be to calculate the average
$$y_m = (y_{n1} + y_{n2} + \dots + y_{ni}) / k$$

Evaluating numeric prediction

.....

- Same strategies: independent test set, cross-validation, significance tests, etc.
- Difference: error measures
- Actual target values: $a_1 \ a_2 \dots a_n$
- Predicted target values: $p_1 \ p_2 \dots p_n$
- Most popular measure: *mean-squared error*

$$\frac{(p_1 - a_1)^2 + \dots + (p_n - a_n)^2}{n}$$

- Easy to manipulate mathematically

See for slides 34-51 this book: Data Mining Practical Machine Learning Tools and Techniques by I. H. Witten, E. Frank, M. A. Hall and C. J. Pal

Other measures

- The *root mean-squared error* :

$$\sqrt{\frac{(p_1 - a_1)^2 + \dots + (p_n - a_n)^2}{n}}$$

- The *mean absolute error* is less sensitive to outliers than the mean-squared error:

$$\frac{|p_1 - a_1| + \dots + |p_n - a_n|}{n}$$

- Sometimes *relative* error values are more appropriate (e.g. 10% for an error of 50 when predicting 500)

Improvement on the mean

.....

- How much does the scheme improve on simply predicting the average?
- The *relative squared error* is:

$$\frac{(p_1 - a_1)^2 + \dots + (p_n - a_n)^2}{(a_1 - \bar{a})^2 + \dots + (a_n - \bar{a})^2}$$

(in this formula and the following two, \bar{a} is the mean value over the training data)

- The *root relative squared error* and the *relative absolute error* are:

$$\sqrt{\frac{(p_1 - a_1)^2 + \dots + (p_n - a_n)^2}{(a_1 - \bar{a})^2 + \dots + (a_n - \bar{a})^2}}$$

$$\frac{|p_1 - a_1| + \dots + |p_n - a_n|}{|a_1 - \bar{a}| + \dots + |a_n - \bar{a}|}$$

Correlation coefficient

.....

- Measures the *statistical correlation* between the predicted values and the actual values

$$\frac{S_{PA}}{\sqrt{S_P S_A}}, \text{ where } S_{PA} = \frac{\sum_i (p_i - \bar{p})(a_i - \bar{a})}{n - 1}, S_P = \frac{\sum_i (p_i - \bar{p})^2}{n - 1},$$
$$S_A = \frac{\sum_i (a_i - \bar{a})^2}{n - 1} \text{ (here, } \bar{a} \text{ is the mean value over the test data)}$$

- Scale independent, between -1 and $+1$
- Good performance leads to large values!

Which measure?

- Best to look at all of them
- Often it doesn't matter
- Example:

Root mean-squared error
Mean absolute error
Root rel squared error
Relative absolute error
Correlation coefficient

A	B	C	D
67.8	91.7	63.3	57.4
41.3	38.5	33.4	29.2
42.2%	57.2%	39.4%	35.8%
43.1%	40.1%	34.8%	30.4%
0.88	0.88	0.89	0.91

- D best, C second-best
 - A, B arguable
-

Regression trees

.....

- Like decision trees, but:
 - Splitting criterion: minimize intra-subset variation
 - Termination criterion: std. dev. becomes small
 - Pruning criterion: based on numeric error measure
 - Prediction: Leaf predicts average class value of instances
- Yields piecewise constant functions
- Easy to interpret
- More sophisticated version: *model trees*

Example Predicting CPU performance

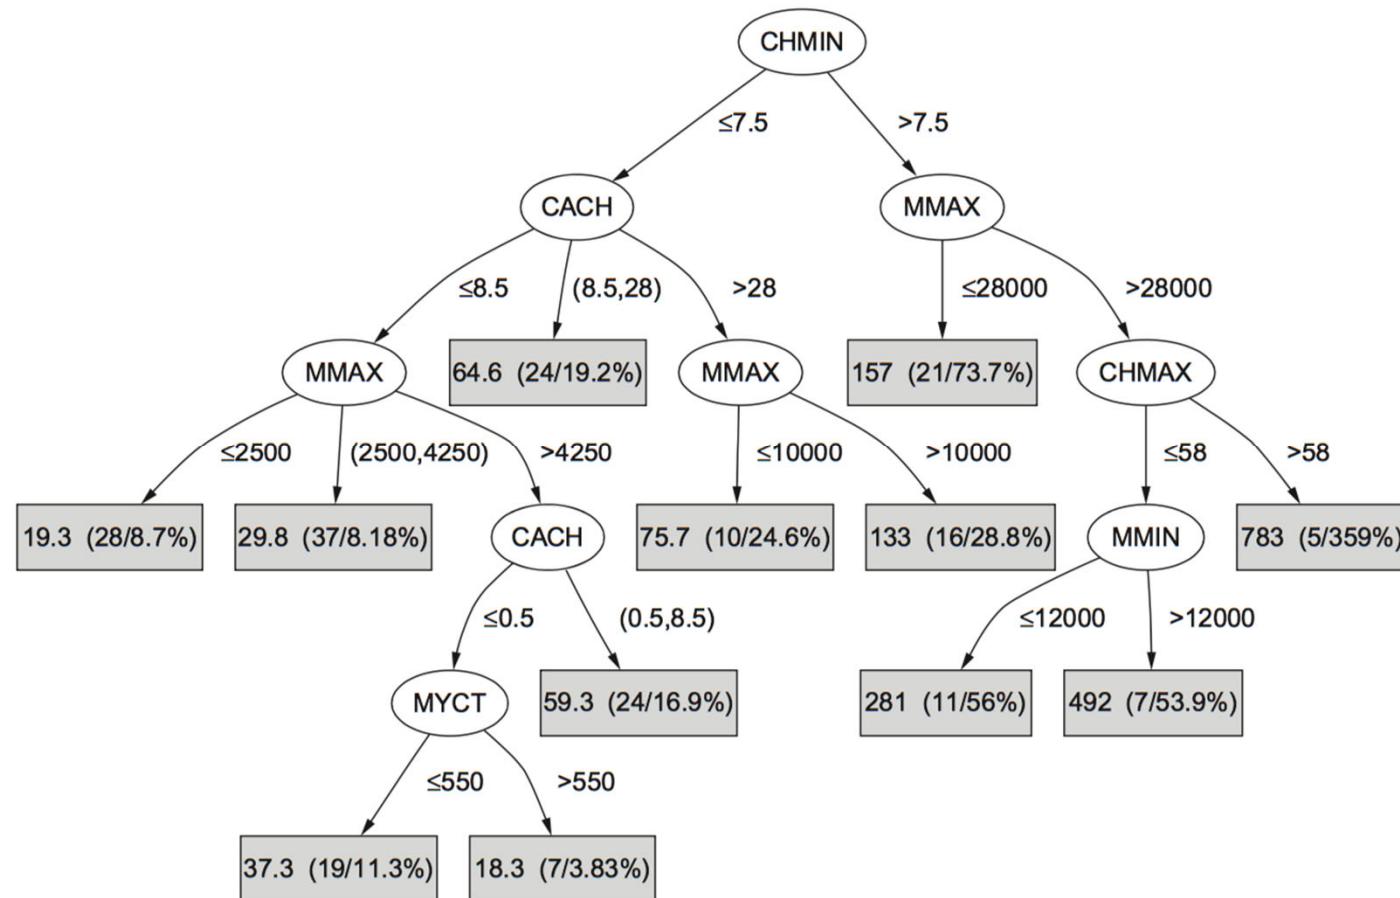
- Example: 209 different computer configurations

	Cycle time (ns)	Main memory (Kb)		Cache (Kb)	Channels		Performance
		MYCT	MMIN	MMAX	CACH	CHMIN	CHMAX
1	125	256	6000	256	16	128	198
2	29	8000	32000	32	8	32	269
...							
208	480	512	8000	32	0	0	67
209	480	1000	4000	0	0	0	45

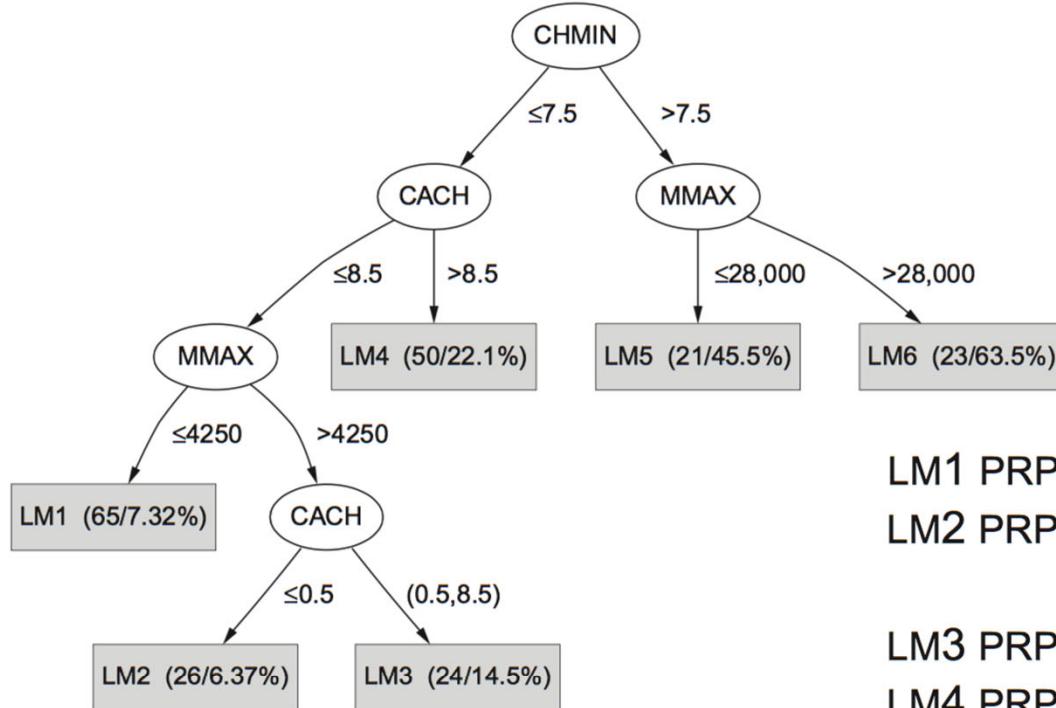
- Linear regression function

$$\begin{aligned} \text{PRP} = & -55.9 + 0.0489 \text{ MYCT} + 0.0153 \text{ MMIN} + 0.0056 \text{ MMAX} \\ & + 0.6410 \text{ CACH} - 0.2700 \text{ CHMIN} + 1.480 \text{ CHMAX} \end{aligned}$$

Regression tree for the CPU data



Model tree for the CPU data



$LM1 \text{ PRP} = 8.29 + 0.004 \text{ MMAX} + 2.77 \text{ CHMIN}$
 $LM2 \text{ PRP} = 20.3 + 0.004 \text{ MMIN} - 3.99 \text{ CHMIN}$
 $+ 0.946 \text{ CHMAX}$
 $LM3 \text{ PRP} = 38.1 + 0.012 \text{ MMIN}$
 $LM4 \text{ PRP} = 19.5 + 0.002 \text{ MMAX} + 0.698 \text{ CACH}$
 $+ 0.969 \text{ CHMAX}$
 $LM5 \text{ PRP} = 285.146 \text{ MYCT} + 1.02 \text{ CACH}$
 $- 9.39 \text{ CHMIN}$
 $LM6 \text{ PRP} = -65.8 + 0.03 \text{ MMIN} - 2.94 \text{ CHMIN}$
 $+ 4.98 \text{ CHMAX}$

Supplementary Material (Optional)

- Build a regression tree
 - Each leaf \Rightarrow linear regression function
 - Smoothing: factor in ancestor's predictions
 - Smoothing formula: $p' = \frac{np + kq}{n + k}$
 - Same effect can be achieved by incorporating ancestor models into the leaves
 - p' \rightarrow the prediction passed up to the next higher node
 - p \rightarrow prediction passed to this node from below
 - q \rightarrow the value predicted by the model at this node
 - n \rightarrow number of training instances that reach the node below
 - k \rightarrow smoothing constant
-

- Build a regression tree
 - Each leaf \Rightarrow linear regression function
 - Smoothing: factor in ancestor's predictions
 - Smoothing formula: $p' = \frac{np + kq}{n + k}$
 - Same effect can be achieved by incorporating ancestor models into the leaves
 - Need linear regression function at each *node*
 - At each node, use only a subset of attributes to build linear regression model
 - Those occurring in subtree
 - (+maybe those occurring in path to the root)
 - Fast: tree usually uses only a small subset of the attributes
-

Building the tree

- Splitting: standard deviation reduction

$$SDR = sd(T) - \sum_i \left| \frac{T_i}{T} \right| \times sd(T_i)$$

- Termination of splitting process:
 - Standard deviation < 5% of its value on full training set
 - Too few instances remain (e.g., < 4)

Pruning:

- Heuristic estimate of absolute error of linear regression models:

$$\frac{n + v}{n - v} \times \text{average_absolute_error}$$

- Greedily remove terms from LR models to minimize estimated error
- Proceed bottom up: compare error of LR model at internal node to error of subtree (this happens before smoothing is applied)
- Heavy pruning: single model may replace whole subtree.....

Nominal attributes

- Convert nominal attributes to binary ones
 - Sort attribute values by their average class values
 - If attribute has k values,
generate $k - 1$ binary attributes
 - i th attribute is 0 if original nominal value is part of the first i nominal values in the sorted list, and 1 otherwise
- Treat binary attributes as numeric in linear regression models and when selecting splits
- Can prove: best SDR split on one of the new binary attributes is the best (binary) SDR split on original nominal attribute
- In practice this process is not applied at every node of the tree but globally at the root node of the tree
 - Splits are no longer optimal but runtime and potential for overfitting are reduced this way.....

Pseudo-code for M5'

.....

- Let us consider the pseudo code for the model tree inducer M5'
 - Four methods:
 - Main method: *MakeModelTree*
 - Method for splitting: *split*
 - Method for pruning: *prune*
 - Method that computes error: *subtreeError*
 - We will briefly look at each method in turn
 - We will assume that the linear regression method performs attribute subset selection based on error (discussed previously)
 - Nominal attributes are replaced globally at the root node
-

```
MakeModelTree (instances)
{
    SD = sd(instances)
    for each k-valued nominal attribute
        convert into k-1 synthetic binary attributes
    root = newNode
    root.instances = instances
    split(root)
    prune(root)
    printTree(root)
}
```

```
split(node)
{
    if sizeof(node.instances) < 4 or
        sd(node.instances) < 0.05*SD
        node.type = LEAF
    else
        node.type = INTERIOR
        for each attribute
            for all possible split positions of attribute
                calculate the attribute's SDR
        node.attribute = attribute with maximum SDR
        split(node.left)
        split(node.right)
}
```

```
prune(node)
{
    if node = INTERIOR then
        prune(node.leftChild)
        prune(node.rightChild)
        node.model = linearRegression(node)
        if subtreeError(node) > error(node) then
            node.type = LEAF
}
```

subtreeError

.....

```
subtreeError(node)
{
    l = node.left; r = node.right
    if node = INTERIOR then
        return (sizeof(l.instances)*subtreeError(l)
                + sizeof(r.instances)*subtreeError(r))
                /sizeof(node.instances)
    else return error(node)
}
```